

CS151 Intro to Data Structures

Hashmaps

Announcements

HW05 due Wednesday

HW06 due next Wednesday (11/22)

Will be released tonight

Lab08 due next Wednesday too

No lab next Wednesday


HW07 due 12/05

Lab09 (today's lab) due then

Notion of a Map

Intuitively, a map M supports the abstraction of using keys as indices with a syntax such as $M[k]$.

Simplest setting is a map with n items using keys that are known to be integers from 0 to $N - 1$, for some $N \geq n$.




More General Keys

What if our keys are not integers in range 0 to $N - 1$?

Use a hash function to map keys to integers into the right range

- Example: last 4 digits of SSN



Hash Functions and Tables

A hash function h maps a key to integers in a fixed interval $[0, N - 1]$

$h(x) = x \% N$ is such a function for integers

$h(x)$ is the *hash value* of key x

A hash table is an array of size N

- associated hash function h
- item (k, v) is stored at index $h(k)$

Example

A hash table storing entries as (SSN, Name), where SSN is a nine-digit positive integer


A hash table is an array of size N

- associated hash function h
- item (k, v) is stored at index $h(k)$

Example

A hash table storing entries as (SSN, Name), where SSN is a nine-digit positive integer

Use an array of size $N = 10000$ and the hash function $h(x) = \text{last 4 digits of } x$



Hash Function

A hash function is usually specified as the composition of two functions:


- hash code: $h_1: \text{key} \rightarrow \text{integers}$
- compression: $h_2: \text{integers} \rightarrow [0, N - 1]$
- $h(x) = h_2(h_1(x))$

The goal is to "disperse" the keys in an appropriately random way


Hash Function Illustration



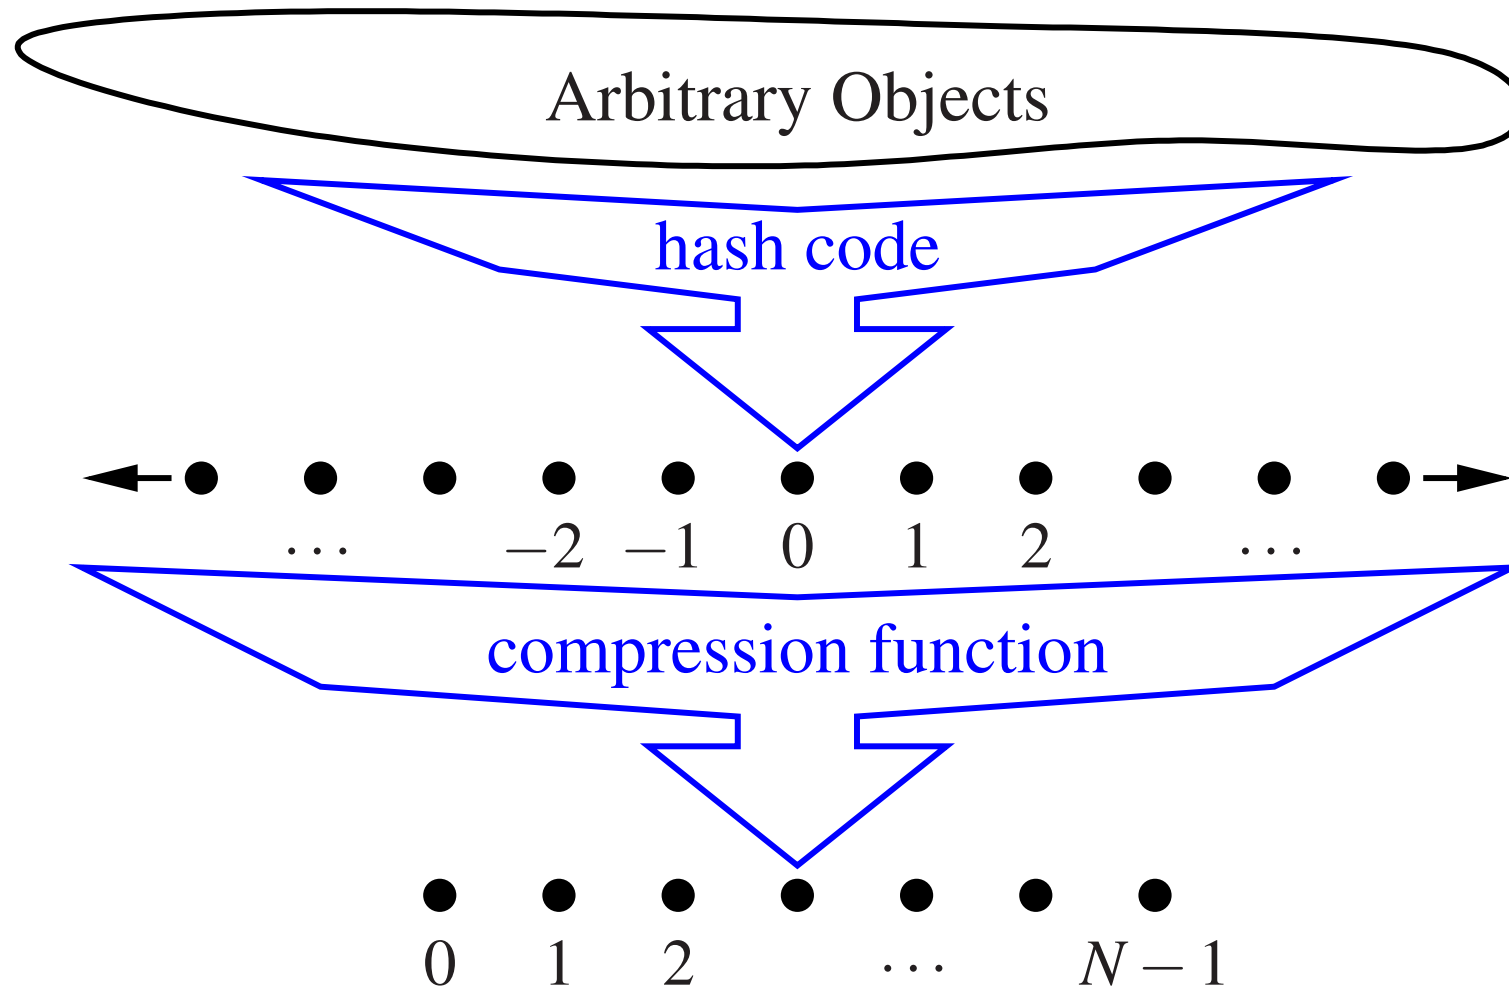
Hash Function Illustration



Hash Function Illustration



Hash Function Illustration



Hash Codes (h_1)

Memory address:

- use the memory address where the keys are stored
- default hash code for Java objects

| | |
|---|--|
| docs.oracle.com/javase/8/docs/api/java/lang/Object.html | |
| Method Summary | |
| All Methods | Instance Methods |
| Concrete Methods | |
| Modifier and Type | Method and Description |
| protected Object | clone() Creates and returns a copy of this object. |
| boolean | equals(Object obj) Indicates whether some other object is "equal to" this on |
| protected void | finalize() Called by the garbage collector on an object when garba object. |
| Class<?> | getClass() Returns the runtime class of this Object. |
| int | hashCode() Returns a hash code value for the object. |

Hash Codes (h_1)

- Memory address:
 - use the memory address where the keys are stored
 - default hash code for Java objects
- Integer cast: interpret the bits storing the keys as integer – `byte`, `short`, `int` and `float`
- Component sum: partition bits into int components and sum them – `long` and `double`

Compression (h_2)

Division: $h_2(x) = x \% N$

N is usually chosen to be a prime

MAD: $h_2(x) = ((ax + b) \% p) \% N$

- N is the size
- p is a prime number, $p > N$
- $0 < a \leq p - 1, 0 \leq b \leq p - 1$
- a scales the range and b shifts the start
- minimize the probability of two keys colliding $-\frac{1}{N}$

AbstractHashMap

```
1 public abstract class AbstractHashMap<K,V> extends AbstractMap<K,V> {
2     protected int n = 0;           // number of entries in the dictionary
3     protected int capacity;        // length of the table
4     private int prime;              // prime factor
5     private long scale, shift;      // the shift and scaling factors
6     public AbstractHashMap(int cap, int p) {
7         prime = p;
8         capacity = cap;
9         Random rand = new Random();
10        scale = rand.nextInt(prime-1) + 1;
11        shift = rand.nextInt(prime);
12        createTable();
13    }
14    public AbstractHashMap(int cap) { this(cap, 109345121); } // default prime
15    public AbstractHashMap() { this(17); } // default capacity
16    // public methods
17    public int size() { return n; }
18    public V get(K key) { return bucketGet(hashValue(key), key); }
19    public V remove(K key) { return bucketRemove(hashValue(key), key); }
20    public V put(K key, V value) {
21        V answer = bucketPut(hashValue(key), key, value);
22        if (n > capacity / 2) // keep load factor <= 0.5
23            resize(2 * capacity - 1); // (or find a nearby prime)
24        return answer;
25    }
26    ... ..
```



AbstractHashMap

```
26 // private utilities
27 private int hashCode(K key) {
28     return (int) ((Math.abs(key.hashCode())*scale + shift) % prime) % capacity);
29 }
30 private void resize(int newCap) {
31     ArrayList<Entry<K,V>> buffer = new ArrayList<>(n);
32     for (Entry<K,V> e : entrySet())
33         buffer.add(e);
34     capacity = newCap;
35     createTable(); // based on updated capacity
36     n = 0; // will be recomputed while reinserting entries
37     for (Entry<K,V> e : buffer)
38         put(e.getKey(), e.getValue());
39 }
40 // protected abstract methods to be implemented by subclasses
41 protected abstract void createTable();
42 protected abstract V bucketGet(int h, K k);
43 protected abstract V bucketPut(int h, K k, V v);
44 protected abstract V bucketRemove(int h, K k);
45 }
```

Collision

A hash function does not guarantee one-to-one mapping – no hash function does

When more than one key hash to the same index, we have a bucket
Each index holds a collection of entries



Collision Handling

Collisions occur when elements with different keys are mapped to the same cell

Separate Chaining: let each cell in the table point to a linked list of entries that map there

Simple, but requires additional memory besides the table

Separate Chaining

using a list-based map at each cell. \mathbb{A} is the table
`get (k) :`

`put (k, v) :`

`remove (k) :`

Skip

skip

Separate Chaining

using a list-based map at each cell. A is the table

```
get(k) :  
    return A[h(k)].get(k)  
put(k, v) :  
    t = A[h(k)].put(k)  
    if t == null then n++    //k is new  
    return t  
remove(k) :  
    t = A[h(k)].remove(k)  
    if t != null then n--    //k found  
    return t
```

ChainHashMap

```
1 public class ChainHashMap<K,V> extends AbstractHashMap<K,V> {
2     // a fixed capacity array of UnsortedTableMap that serve as buckets
3     private UnsortedTableMap<K,V>[] table; // initialized within createTable
4     public ChainHashMap() { super(); }
5     public ChainHashMap(int cap) { super(cap); }
6     public ChainHashMap(int cap, int p) { super(cap, p); }
7     /** Creates an empty table having length equal to current capacity. */
8     protected void createTable() {
9         table = (UnsortedTableMap<K,V>[]) new UnsortedTableMap[capacity];
10    }
11    /** Returns value associated with key k in bucket with hash value h, or else null. */
12    protected V bucketGet(int h, K k) {
13        UnsortedTableMap<K,V> bucket = table[h];
14        if (bucket == null) return null;
15        return bucket.get(k);
16    }
17    /** Associates key k with value v in bucket with hash value h; returns old value. */
18    protected V bucketPut(int h, K k, V v) {
19        UnsortedTableMap<K,V> bucket = table[h];
20        if (bucket == null)
21            bucket = table[h] = new UnsortedTableMap<>();
22        int oldSize = bucket.size();
23        V answer = bucket.put(k,v);
24        n += (bucket.size() - oldSize); // size may have increased
25        return answer;
26    }
```


ChainHashMap

```
27  /** Removes entry having key k from bucket with hash value h (if any). */
28  protected V bucketRemove(int h, K k) {
29      UnsortedTableMap<K,V> bucket = table[h];
30      if (bucket == null) return null;
31      int oldSize = bucket.size();
32      V answer = bucket.remove(k);
33      n -= (oldSize - bucket.size());    // size may have decreased
34      return answer;
35  }
36  /** Returns an iterable collection of all key-value entries of the map. */
37  public Iterable<Entry<K,V>> entrySet() {
38      ArrayList<Entry<K,V>> buffer = new ArrayList<>();
39      for (int h=0; h < capacity; h++)
40          if (table[h] != null)
41              for (Entry<K,V> entry : table[h].entrySet())
42                  buffer.add(entry);
43      return buffer;
44  }
45 }
```

Separate Chaining

using a list-based map at each cell. A is the table

```
get(k) :  
    return A[h(k)].get(k)  
put(k, v) :  
    t = A[h(k)].put(k)  
    if t == null then n++    //k is new  
    return t  
remove(k) :  
    t = A[h(k)].remove(k)  
    if t != null then n--    //k found  
    return t
```

Open Addressing and Probing


- Colliding item is put in a different cell
- Linear probing: place the colliding item in the next (circularly) available table cell
- Colliding items cluster together – future collisions to cause a longer sequence of probes
- Example: $h(x) = x \% 13$
 - insert 18(5), 41(2), 22(9), 44(5), 59(7), 32(6), 31(5), 73(8)

skip


Skip

Open Addressing and Probing

- Example: $h(x) = x \% 13$
- insert 18(5), 41(2), 22(9), 44(5), 59(7), 32(6), 31(5), 73(8)



- collision: 4



ProbeHashMap

```
1 public class ProbeHashMap<K,V> extends AbstractHashMap<K,V> {
2     private MapEntry<K,V>[] table;           // a fixed array of entries (all initially null)
3     private MapEntry<K,V> DEFUNCT = new MapEntry<>(null, null); //sentinel
4     public ProbeHashMap() { super(); }
5     public ProbeHashMap(int cap) { super(cap); }
6     public ProbeHashMap(int cap, int p) { super(cap, p); }
7     /** Creates an empty table having length equal to current capacity. */
8     protected void createTable() {
9         table = (MapEntry<K,V>[]) new MapEntry[capacity]; // safe cast
10    }
11    /** Returns true if location is either empty or the "defunct" sentinel. */
12    private boolean isAvailable(int j) {
13        return (table[j] == null || table[j] == DEFUNCT);
14    }
```

ProbeHashMap

```
15  /** Returns index with key k, or  $-(a+1)$  such that k could be added at index a. */
16  private int findSlot(int h, K k) {
17      int avail = -1; // no slot available (thus far)
18      int j = h; // index while scanning table
19      do {
20          if (isAvailable(j)) { // may be either empty or defunct
21              if (avail == -1) avail = j; // this is the first available slot!
22              if (table[j] == null) break; // if empty, search fails immediately
23          } else if (table[j].getKey().equals(k))
24              return j; // successful match
25          j = (j+1) % capacity; // keep looking (cyclically)
26      } while (j != h); // stop if we return to the start
27      return  $-(avail + 1)$ ; // search has failed
28  }
29  /** Returns value associated with key k in bucket with hash value h, or else null. */
30  protected V bucketGet(int h, K k) {
31      int j = findSlot(h, k);
32      if (j < 0) return null; // no match found
33      return table[j].getValue();
34  }
```


ProbeHashMap

```
35  /** Associates key k with value v in bucket with hash value h; returns old value. */
36  protected V bucketPut(int h, K k, V v) {
37      int j = findSlot(h, k);
38      if (j >= 0)                                // this key has an existing entry
39          return table[j].setValue(v);
40      table[-(j+1)] = new MapEntry<>(k, v);    // convert to proper index
41      n++;
42      return null;
43  }
44  /** Removes entry having key k from bucket with hash value h (if any). */
45  protected V bucketRemove(int h, K k) {
46      int j = findSlot(h, k);
47      if (j < 0) return null;                    // nothing to remove
48      V answer = table[j].getValue();
49      table[j] = DEFUNCT;                      // mark this slot as deactivated
50      n--;
51      return answer;
52  }
53  /** Returns an iterable collection of all key-value entries of the map. */
54  public Iterable<Entry<K,V>> entrySet() {
55      ArrayList<Entry<K,V>> buffer = new ArrayList<>();
56      for (int h=0; h < capacity; h++)
57          if (!isAvailable(h)) buffer.add(table[h]);
58      return buffer;
59  }
60 }
```

Probing Distance


- Given a hash value $h(x)$, linear probing generates $h(x), h(x) + 1, h(x) + 2, \dots$
- Primary clustering – the bigger the cluster gets, the faster it grows
- **Quadratic probing** – $h(x), h(x) + 1, h(x) + 4, h(x) + 9, \dots$
- Quadratic probing leads to secondary clustering, more subtle, not as dramatic, but still systematic

Double Hashing

- Interval between probes is fixed but computed by a second hash function
- Use a secondary hash function $d(k)$ to handle collisions by placing an item in the first available cell of the series
$$(h(k) + i \times d(k)) \% N, \quad 0 \leq i \leq N - 1$$
- N must be prime
- $d(k) = q - k \% q, q < N, q$ is prime

Example

- Double hashing:
 - $N = 13$
 - $h(k) = k \% 13$
 - $d(k) = 7 - k \% 7$
- Insert 18, 41, 22, 44, 59, 32, 31, 73




Skip slide

Skip slide

Example

- Double hashing:
 - $N = 13$
 - $h(k) = k \% 13$
 - $d(k) = 7 - k \% 7$
- Insert 18, 41, 22, 44, 59, 32, 31, 73
- Collision: 2

| k | $h(k)$ | $d(k)$ | Probes | |
|-----|--------|--------|--------|-----|
| 18 | 5 | 3 | 5 | |
| 41 | 2 | 1 | 2 | |
| 22 | 9 | 6 | 9 | |
| 44 | 5 | 5 | 5 | 10 |
| 59 | 7 | 4 | 7 | |
| 32 | 6 | 3 | 6 | |
| 31 | 5 | 4 | 5 | 9 0 |
| 73 | 8 | 4 | 8 | |



Performance Analysis

- In the worst case, searches, insertions and removals take $O(n)$ time
 - when all the keys collide
- The load factor $\alpha = n/N$ affects the performance of a hash table
 - expected number of probes for an insertion with open addressing is $\frac{1}{1-\alpha}$
- Expected time of all operations is $O(1)$ provided α is not close to 100%

Open Addressing vs Chaining

- Probing is significantly faster in practice
- locality of references – much faster to access a series of elements in an array than to follow the same number of pointers in a linked list
- Efficient probing requires soft/lazy deletions – tombstoning, why?
- May require graveyard defragmenting

Probing Tradeoffs

- Linear probing – best cache performance but most sensitive to clustering
- Double hashing – poor cache performance but exhibits virtually no clustering
- Quadratic – inbetween
- As load factor approaches 100%, number of probes rises dramatically
- Even with good hash functions, keep load factor 80% or below (50% is typical)
- Other open addressing methods besides probing

Good Hash Function

- is critical to performance
- A poor hash function can lead to poor performance even at very low load factor
- It is easy to unintentionally write a hash function that leads to severe clustering
- Testing your hash function is paramount
- stick with `.hashCode()`

Implementation

- Interface `Entry` defines the expected behaviors of an entry (key-value pair)
- Interface `Map` defines the expected behaviors of a hashmap as an ADT

Implementation

- Abstract class `AbstractMap` implements `Map`
 - base class to provide support for misc utilities
 - `isEmpty`, `toString`
 - nested class `MapEntry` implements `Entry`
 - iterators
 - `keySet`, `values`
 - depends on `entrySet` to be provided by concrete child class

Implementation

- `AbstractHashMap` **extends** `AbstractMap`
 - **base class to provide support for hashing**
 - `capacity`, `prime`, `scale`, `shift`, `n`
 - **expands table when out of space on put**
 - `.hashCode` implements h_1
 - `hashValue` implements h_2 with MAD

Implementation

- `ProbeHashMap` **extends** `AbstractHashMap`
 - Concrete class
 - Open-addressing with linear probing
 - `bucketGet`, `bucketPut`, `bucketRemove`
 - `findSlot`

Performance of Hashtable

| | Hash Expected | Hash Worst |
|---------|---------------|------------|
| search | | |
| insert | | |
| remove | | |
| min/max | | |

| | Unsorted array | Sorted array | Unsorted list | Sorted list | BST balanced | Hash Expected |
|---------|----------------|--------------|---------------|-------------|--------------|---------------|
| search | $O(n)^*$ | $O(\log n)$ | $O(n)$ | $O(n)$ | $O(\log n)$ | $O(1)$ |
| insert | $O(1)^*$ | $O(n)$ | $O(1)$ | $O(n)$ | $O(\log n)$ | $O(1)$ |
| remove | $O(1)^*$ | $O(n)$ | $O(1)$ | $O(1)$ | $O(\log n)$ | $O(1)$ |
| min/max | $O(n)$ | $O(1)$ | $O(n)$ | $O(1)$ | $O(\log n)$ | $O(n)$ |

Performance of Hashtable

| | Hash Expected | Hash Worst |
|---------|---------------|------------|
| search | $O(1)$ | $O(n)$ |
| insert | $O(1)$ | $O(n)$ |
| remove | $O(1)$ | $O(n)$ |
| min/max | $O(n)$ | $O(n)$ |

| | Unsorted array | Sorted array | Unsorted list | Sorted list | BST balanced | Hash Expected |
|---------|----------------|--------------|---------------|-------------|--------------|---------------|
| search | $O(n)^*$ | $O(\log n)$ | $O(n)$ | $O(n)$ | $O(\log n)$ | $O(1)$ |
| insert | $O(1)^*$ | $O(n)$ | $O(1)$ | $O(n)$ | $O(\log n)$ | $O(1)$ |
| remove | $O(1)^*$ | $O(n)$ | $O(1)$ | $O(1)$ | $O(\log n)$ | $O(1)$ |
| min/max | $O(n)$ | $O(1)$ | $O(n)$ | $O(1)$ | $O(\log n)$ | $O(n)$ |

Hashtable vs Array

- A hashtable is an unsorted array with a fast search – $O(1)$ expected
- An array is more memory efficient, but slower for searching (without key-index pairing)
- If your data has natural indexing – a way to assign/associate an ID/unique integer to each entry, then you are better off using an array. You have a hash function with 1-to-1 mapping and guaranteed no collisions

Hashtable Size

- Should be a prime
- twice the size of max number of keys
- or 1.3 times if n is very large
- $1/1.333 = 75\%$ load factor
- Keep track of load factor and expand (rehash) the hash table when necessary