

CS151 Intro to Data Structures

Sorted Maps, Tree Maps,
Balanced Search Trees, AVL Trees

Announcements

HW06 due Wednesday 11/29
Lab08 due Wednesday too

No lab this week

Last lab will be next week

HW07 due Wednesday 12/06
Hashmaps & Sorting


HW08 due 12/14

Outline


Sorting Review

Sorted Maps

Sorting



Sorting



Step: 1


In-place Partitioning (Hoare's)

- Use two indices to split into L and EUG_r

6 2 5 1 0 7 3 5 9 2 7 9 8 9 7 9 3


- Repeat until l and r cross:
 - Move l to the right to find $\geq x$
 - Move r to the left to find $< x$
 - swap elements at l and r

pivot = 6
use first
element as
pivot or
swap pivot
with last




Example with Duplicates

- Use last element in array as pivot





Example with Duplicates

- Use last element in array as pivot




Example with Duplicates

- Use last element in array as pivot




Example with Duplicates

- Use last element in array as pivot




Example with Duplicates

- Use last element in array as pivot





Example with Duplicates

- Use last element in array as pivot




Example with Duplicates

- Use last element in array as pivot



HashMap


Keys are uniformly distributed in a hashtable



HashMap

Keys are uniformly distributed in a hashtable

Key: Timestamp; Value: Action




HashMap

Keys are uniformly distributed in a hashtable

Key: Timestamp; Value: Action

A: (10:14, Deposit)




HashMap

Keys are uniformly distributed in a hashtable

Key: Timestamp; Value: Action

$A = (10:14, \text{Deposit})$

$P(h(A) = 0)$




HashMap

Keys are uniformly distributed in a hashtable

Key: Timestamp; Value: Action

$A = (10:14, \text{Deposit})$

$$P(h(A) = 0) = 1/10$$



HashMap

Find all transactions within a time range?

Find all deposits?

(10:14, Deposit)	(1:59, Withdraw)			(10:15, Withdraw)			(2:04, Deposit)	(10:16, Transfer)	
---------------------	---------------------	--	--	----------------------	--	--	--------------------	----------------------	--

Sorted Map ADT

Similar items are stored near each other in clusters

`firstEntry()`: Returns the entry with smallest key value (or null, if the map is empty).

`lastEntry()`: Returns the entry with largest key value (or null, if the map is empty).

`ceilingEntry(k)`: Returns the entry with the least key value greater than or equal to k (or null, if no such entry exists).

`floorEntry(k)`: Returns the entry with the greatest key value less than or equal to k (or null, if no such entry exists).

`lowerEntry(k)`: Returns the entry with the greatest key value strictly less than k (or null, if no such entry exists).

`higherEntry(k)`: Returns the entry with the least key value strictly greater than k (or null if no such entry exists).

`subMap(k_1 , k_2)`: Returns an iteration of all entries with key greater than or equal to k_1 , but strictly less than k_2 .

Sorted Maps

Used on data that has a natural ordering – comparable elements, timed events

Implement with a sorted array - relies on binary search

Fast search for min/max key, key and nearby keys - $O(\log n)$

Slow updates - $O(n)$

`java.util.NavigableMap`

`java.util.SortedMap`

Sorted Map Analysis

Method	Running Time
size	$O(1)$
get	$O(\log n)$
put	$O(n)$; $O(\log n)$ if map has entry with given key
remove	$O(n)$
firstEntry, lastEntry	$O(1)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(\log n)$
subMap	$O(s + \log n)$ where s items are reported
entrySet, keySet, values	$O(n)$

TreeMap

A sorted map implemented with a binary search tree

Recall the fundamental methods of a map:

`get(k)`: Returns the value *v* associated with key *k*, if such an entry exists; otherwise returns null.

`put(k, v)`: Associates value *v* with key *k*, replacing and returning any existing value if the map already contains an entry with key equal to *k*.

`remove(k)`: Removes the entry with key equal to *k*, if one exists, and returns its value; otherwise returns null.


Binary Search Trees

Performance is directly affected by the height of tree

All operations are $O(h)$

- $h = O(n)$ worst case
- $h = O(\log n)$ best case

Expected $O(\log n)$ if tree is balanced




Balanced Search Trees

A variety of algorithms that augments a standard BST with occasional operations to reshape and reduce height

Rotation:

- move a child to be above its parent and relink subtrees to maintain BST order
- $O(1)$



Tree Rotation

Rotation can be to the right or left


Rotate reduces/increases the depth of nodes in subtrees T_1 and T_3 by 1

Rotation maintains BST order


Rotate is $O(1)$

One or more rotations can be combined to provide broader rebalancing


Tri-node restructuring: a node x , its parent y and its grandparent z




Single Rotation (around z)




Single Rotation (around z)



Single Rotation (around z)



Single Rotation (around z)




Rotations

Right rotation:



- Root node's left child becomes the new root
- Root node becomes the left child's right child

Left rotation:




- Root node's right child becomes the new root
- Root node becomes the right child's left child




Rotation




Rotation




Double Rotation




Double Rotation




Double Rotation (around z)



Double Rotation (around z)



Double Rotation (around z)

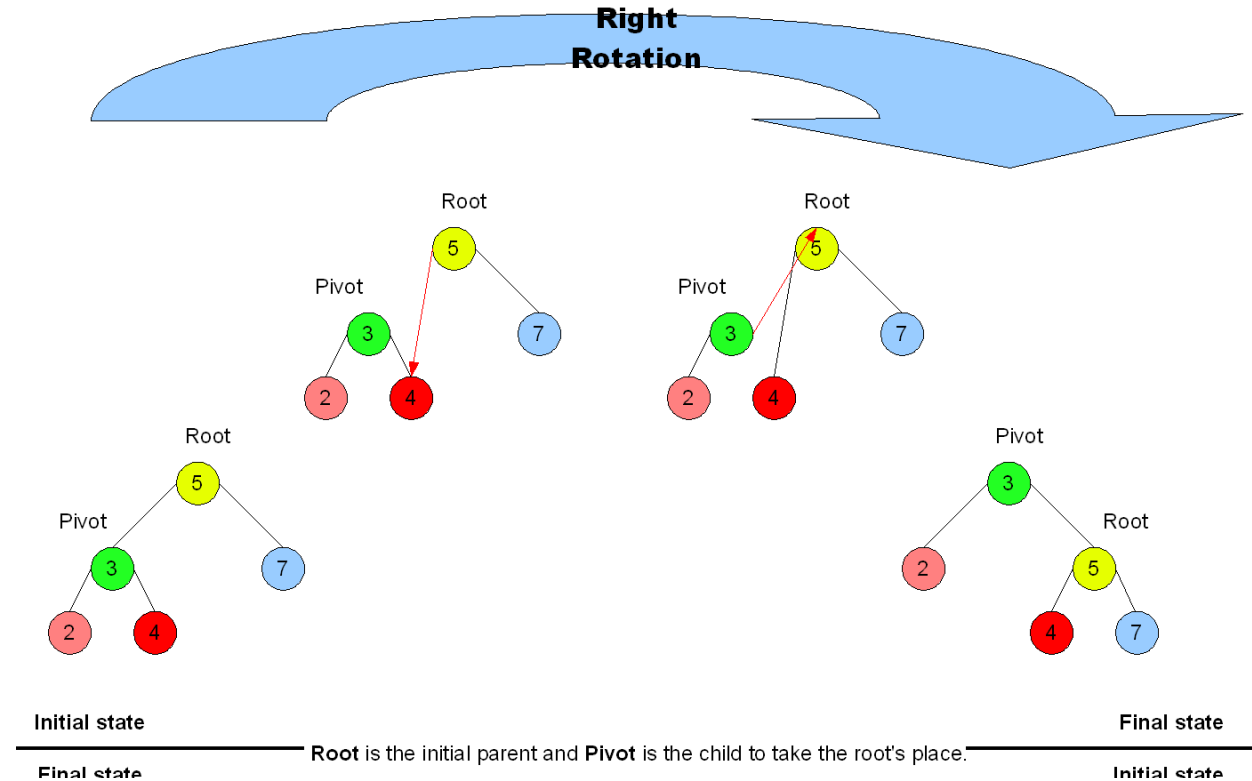


Tree Rotations

```

rotateRight(r):
  if (r.left==null) return
  p = r.left
  r.left = p.right
  p.right = r

  // set parent
  if r.parent == null
    root = p
    p.parent = null
  else
    if(r.parent.left == r)
      r.parent.left=p
    else
      r.parent.right=p
  
```





Initial state

Final state

Final state

Root is the initial parent and **Pivot** is the child to take the root's place.

Initial state



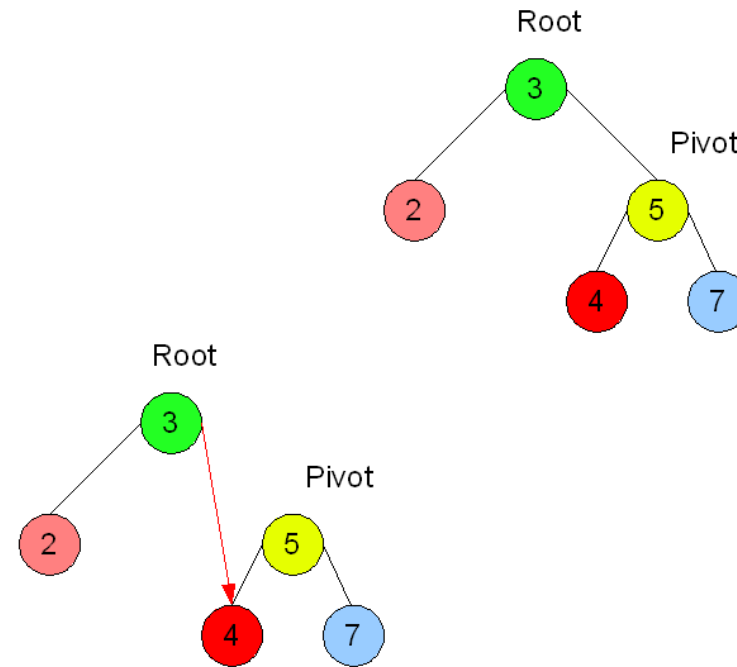
Initial state

Final state

Final state

Root is the initial parent and **Pivot** is the child to take the root's place.

Initial state



**Left
Rotation**

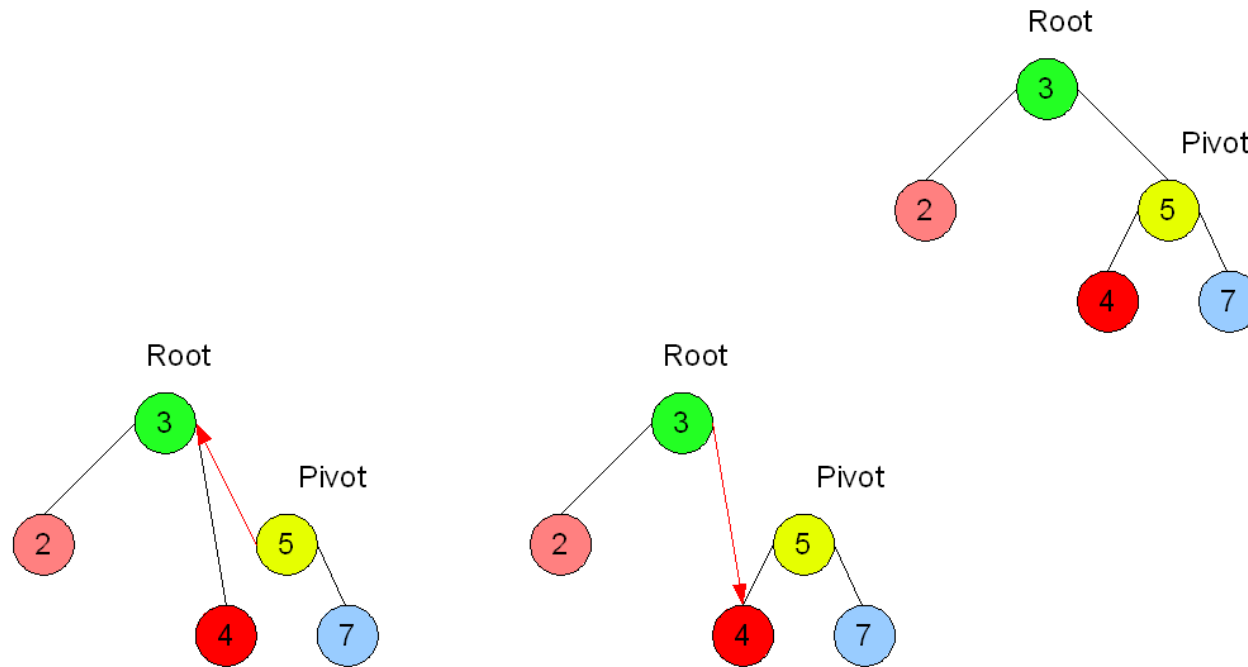
Initial state

Final state

Final state

Root is the initial parent and **Pivot** is the child to take the root's place.

Initial state



**Left
Rotation**






Initial state

Final state

Final state

Root is the initial parent and **Pivot** is the child to take the root's place.

Initial state



AVL Tree

Height of a subtree is the number of edges on the longest path from subtree root to a leaf


Height-balance property

- For every internal node, the heights of the two children differ by at most 1

Any binary tree satisfying the height-balance property is an AVL tree

AVL Tree Example

- leaves are sentinels and have height 0




AVL height

The height of an AVL is $O(\log n)$

$n(h)$ denotes the number of minimum internal nodes for an AVL with height h

- $n(1) = 1$ and $n(2) = 2$
- $n(h) = 1 + n(h-1) + n(h-2)$
- $n(h) > 2 \cdot n(h-2) > 2^i \cdot n(h-2i)$
- $h - 2i = 1 \implies i = \frac{h}{2} - 1$
- $\log(n(h)) = \frac{h}{2} - 1 \implies h < 2 \log(n(h)) + 1$

Insert 54




Insertion (54)

New node always has height 1



Parent may change height

All ancestors may become unbalanced

Perform rotations for unbalanced ancestors




$O(1)$ Rotation Restores Global Balance



Insert into T_3


After rebalance:

- x , y and z are balanced after
- root of subtree returns to height $h + 2$, as before



Exercise

- Create an AVL tree by inserting the nodes in this order:
 - M, N, O, L, K, Q, P, H, I, A




- AVL balance marked on nodes
- $\text{balance}(n) = \text{height of right subtree} - \text{height of left subtree}$
- AVL balance property: $|\text{balance}(n)| \leq 1$

AVL Animation



Exercise

- Create an AVL tree by inserting the nodes in this order:
 - M, N, O, L, K, Q, P, H, I, A



- AVL balance marked on nodes
- $\text{balance}(n) = \text{height of right subtree} - \text{height of left subtree}$
- AVL balance property: $|\text{balance}(n)| \leq 1$

Rebalance: no null checks

```
rebalance(n):  
    updateHeight(n) // update height from children  
    lh = n.left.height rh = n.right.height  
    if (lh > rh+1) // left subtree too tall  
        llh = n.left.left.height lrh = n.left.right.height  
        if (llh >= lrh)  
            return rotateRight(n) //left-left  
        else  
            return rotateLeftRight(n) //left-right  
    else if (rh > lh+1) // right subtree too tall  
        // ... symmetric  
    else return n // no rotation
```

Helpers

```
rotateRight(r):  
    p = r.left  
    r.left = p.right  
    p.right = r  
    updateHeight(r)  
    updateHeight(p)  
    // let caller set parent  
    // return new subtree root  
    return p
```


```
rotateLeftRight(r):  
    r.left = rotateLeft(r.left)  
    return rotateRight(r)
```

```
updateHeight(n):  
    lh = n.left.height  
    rh = n.right.height  
    height = 1+max(lh, rh)
```

Insert with parent

```
insertRec(root, key):  
    if root == null:  
        return new Node(key)  
    if root.key > key:  
        root.left = insertRec(root.left, key)  
        root.left.parent = root  
    else  
        root.right = insertRec(root.right, key)  
        root.right.parent = root  
    return root
```


Delete 32



Deletion


Deletion structurally removes a node with 0 or 1 child

- predecessor has 0 or 1 left child
- successor has 0 or 1 right child

Deletion may reduce the height of parent

Ancestors may become unbalanced

Rotate to rebalance just like insertion



$O(\log n)$ Rotations

Unlike insertion where rotation of the nearest unbalanced ancestor restores the balance globally

On deletion, rotation of the nearest unbalanced ancestor only guarantees balance locally to the subtree

Worst-case requires $O(\log n)$ rotations up the tree to restore balance globally

Performance of AVLTreeMap

Method	Running Time
size, isEmpty	$O(1)$
get, put, remove	$O(\log n)$
firstEntry, lastEntry	$O(\log n)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(\log n)$
subMap	$O(s + \log n)$
entrySet, keySet, values	$O(n)$

Book's Implementation of AVL

- 17 classes!

- Interfaces

- Entry
- Position
- Queue
- Tree
- BinaryTree
- Map
- SortedMap

- Abstract classes:

- AbstractTree
- AbstractBinaryTree
- AbstractMap
- AbstractSortedMap

- Concrete classes

- SinglyLinkedList
- LinkedQueue
- LinkedBinaryTree
- TreeMap
- AVLTreeMap