

CS151 Intro to Data Structures

Balanced Search Trees, AVL Trees

Announcements

Last lab will be this week

HW07 due Friday 12/08
Hashmaps & Sorting

HW08 due 12/14

HW07

- Check correctness on small `ArrayList<Integer>`
- All five of your deduplication methods should return identical lists
- The one that calls `Collections.sort` is probably most trustworthy
- `doubleHash` should result in better hashing stats over `linearHash`

Faculty Interview/Mock Lecture

- Friday 12/08 – 11-11am
- Binary Search Tree
- Location: TBD
- Tea & Snacks

Outline

Review Balanced Binary Trees

AVL Trees

Splay Trees

Skip List

Sets


Binary Search Trees

Performance is directly affected by the height of tree

All operations are $O(h)$

- $h = O(n)$ worst case
- $h = O(\log n)$ best case

Expected $O(\log n)$ if tree is balanced



Balanced Trees





- The difference between the height of the left and right subtree for any node is at most 1
- Left subtree of a node is balanced
- Right subtree of a node is balanced

Types of imbalances

- Left Left
- Left Right
- Right Left
- Right Right

Types of imbalances

- Left Left
- Left Right
- Right Left
- Right Right




Rotations

Right rotation:

- Root node's left child becomes the new root
- Root node becomes the left child's right child





Left rotation:

- Root node's right child becomes the new root
- Root node becomes the right child's left child



Types of imbalances

- Left Left
 - Right rotation
- Left Right
 - Left rotation (makes LL)
 - Then Right rotation
- Right Left
 - Right rotation (makes RR)
 - Then left rotation
- Right Right
 - Left rotation



Outline

Double Hashing Review (Homework 07)

Balanced Binary Trees

AVL Trees

Splay Trees

Red-Black Trees

AVL Tree

Height of a subtree is the number of edges on the longest path from subtree root to a leaf


Height-balance property

- For every internal node, the heights of the two children differ by at most 1

Any binary tree satisfying the height-balance property is an AVL tree


AVL Tree Example

- leaves are sentinels and have height 0



AVL Tree Example

- leaves are sentinels and have height 0




AVL height

The height of an AVL is $O(\log n)$

$n(h)$ denotes the number of minimum internal nodes for an AVL with height h

- $n(1) = 1$ and $n(2) = 2$
- $n(h) = 1 + n(h-1) + n(h-2)$
- $n(h) > 2 \cdot n(h-2) > 2^i \cdot n(h-2i)$
- $h - 2i = 1 \implies i = \frac{h}{2} - 1$
- $\log(n(h)) = \frac{h}{2} - 1 \implies h < 2 \log(n(h)) + 1$

Insert 54




Insertion (54)

New node always has height 1



Parent may change height

All ancestors may become unbalanced

Perform rotations for unbalanced ancestors




$O(1)$ Rotation Restores Global Balance



Insert into T_3


After rebalance:

- x , y and z are balanced after
- root of subtree returns to height $h + 2$, as before



Exercise

- Create an AVL tree by inserting the nodes in this order:
 - M, N, O, L, K, Q, P, H, I, A




- AVL balance marked on nodes
- $\text{balance}(n) = \text{height of right subtree} - \text{height of left subtree}$
- AVL balance property: $|\text{balance}(n)| \leq 1$

AVL Animation



Exercise

- Create an AVL tree by inserting the nodes in this order:
 - M, N, O, L, K, Q, P, H, I, A



- AVL balance marked on nodes
- $\text{balance}(n) = \text{height of right subtree} - \text{height of left subtree}$
- AVL balance property: $|\text{balance}(n)| \leq 1$

Rebalance: no null checks

```
rebalance(n):  
    updateHeight(n) // update height from children  
    lh = n.left.height rh = n.right.height  
    if (lh > rh+1) // left subtree too tall  
        llh = n.left.left.height lrh = n.left.right.height  
        if (llh >= lrh)  
            return rotateRight(n) //left-left  
        else  
            return rotateLeftRight(n) //left-right  
    else if (rh > lh+1) // right subtree too tall  
        // ... symmetric  
    else return n // no rotation
```

Helpers

```
rotateRight(r):  
    p = r.left  
    r.left = p.right  
    p.right = r  
    updateHeight(r)  
    updateHeight(p)  
    // let caller set parent  
    // return new subtree root  
    return p
```


```
rotateLeftRight(r):  
    r.left = rotateLeft(r.left)  
    return rotateRight(r)
```

```
updateHeight(n):  
    lh = n.left.height  
    rh = n.right.height  
    height = 1+max(lh, rh)
```


Insert with parent

```
insertRec(root, key):  
    if root == null:  
        return new Node(key)  
    if root.key > key:  
        root.left = insertRec(root.left, key)  
        root.left.parent = root  
    else  
        root.right = insertRec(root.right, key)  
        root.right.parent = root  
    return root
```

Delete 32



Deletion

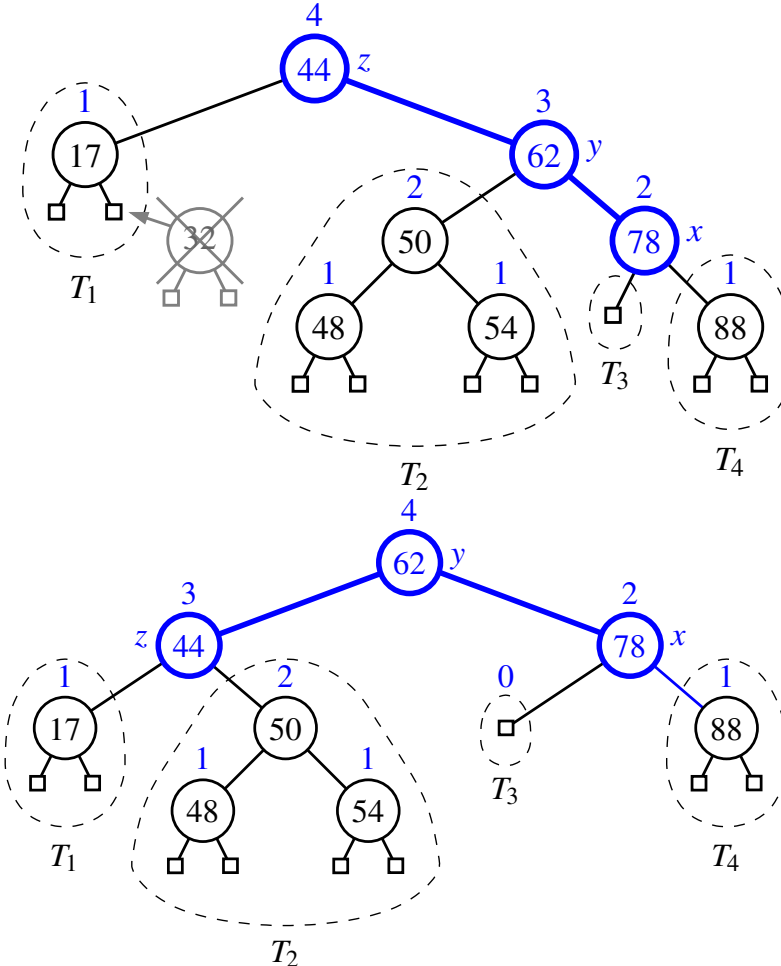
Deletion structurally removes a node with 0 or 1 child

- predecessor has 0 or 1 left child
- successor has 0 or 1 right child

Deletion may reduce the height of parent

Ancestors may become unbalanced

Rotate to rebalance just like insertion



$O(\log n)$ Rotations

Unlike insertion where rotation of the nearest unbalanced ancestor restores the balance globally

On deletion, rotation of the nearest unbalanced ancestor only guarantees balance locally to the subtree

Worst-case requires $O(\log n)$ rotations up the tree to restore balance globally

Performance of AVLTreeMap

Method	Running Time
size, isEmpty	$O(1)$
get, put, remove	$O(\log n)$
firstEntry, lastEntry	$O(\log n)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(\log n)$
subMap	$O(s + \log n)$
entrySet, keySet, values	$O(n)$

Book's Implementation of AVL

- 17 classes!

- Interfaces

- Entry
- Position
- Queue
- Tree
- BinaryTree
- Map
- SortedMap

- Abstract classes:

- AbstractTree
- AbstractBinaryTree
- AbstractMap
- AbstractSortedMap

- Concrete classes

- SinglyLinkedList
- LinkedQueue
- LinkedBinaryTree
- TreeMap
- AVLTreeMap

Outline

Review Balanced Binary Trees

AVL Trees

Splay Trees

Skip List


Sets

Splay Tree

- A binary search tree that doesn't enforce a $O(\log n)$ bound on the height
- Efficiency is achieved due to a move-to-root operation, called splaying
- Performed at the leaf reached during every insert, delete and search
- Causes the more frequently accessed elements to be near the top


Splaying

- Swapping a BST node x up depends on the relative position of x , its parent y and its grandparent z
- Zig/zag: y has no parent
- Splaying will continue these rotations until x becomes root




Splaying


- zig-zig (zag-zag):
 x and y are both
right/left children




- zig-zag (zag-zig):
one right one left



Example – insert 14




Example




When/what to Splay

- On search for x : if x is found, splay x else splay x 's parent
- On insert x : splay x after insertion

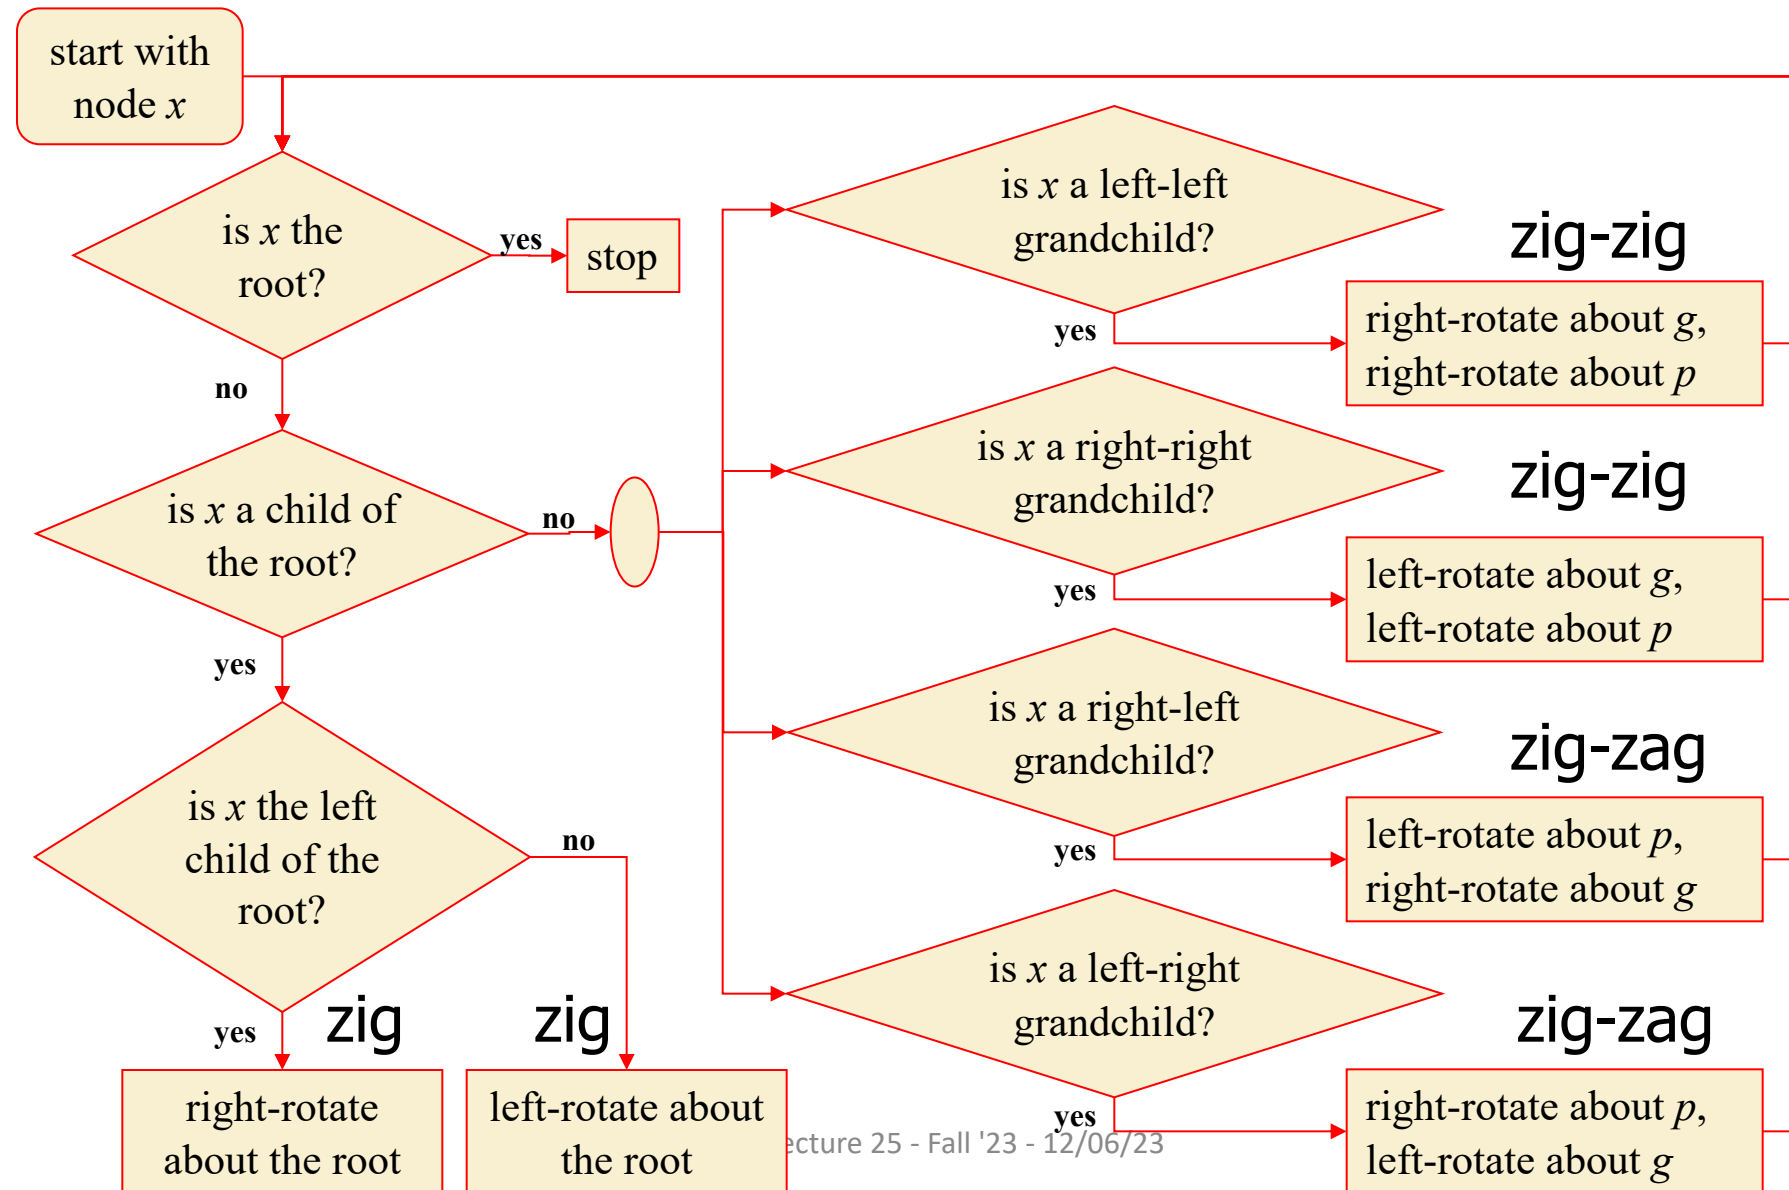


- On delete x : splay parent of removed node
 - x is removed
 - in-order successor/predecessor removed

Deletion



How to Splay



Analysis of Splaying

- Splay trees do rotations after every operation (even search)
- Runtime of each search/insert/delete is proportional to the time for splaying
- Each zig-zig, zig-zag or zig is $O(1)$
- Splaying a node at height h is $O(h)$
- Worst case height of a splay tree is $O(n)$

Amortized Performance

- A splay tree performs well in amortization – in a sequence of mixed searches, insertions and deletions
- Splay tree performs better for many sequences of non-random operations
- Amortized cost for any splay operation is $O(\log n)$
- Must faster search than $O(\log n)$ on frequently requested items

AVL Rotations

- AVL insert – $O(\log n)$
 - Find the lowest out-of-balance ancestor – also known as the critical node, rotate critical node to balance. Loop ends after single rotation
 - $O(\log n)$ search up the tree to find critical node + $O(1)$ rotations
- AVL delete - $O(\log n)$
 - $O(\log n)$ rotations on delete

Outline

Review Balanced Binary Trees

AVL Trees

Splay Trees

Skip List

Sets


The Problem with lists

- If you must use a linked list
 - because of frequent insertions and deletions
- How do you arrange a fast search?
- What if the list is sorted?
- Still no way to arrange binary search
- `java.util.Concurrent.ConcurrentSkipListMap`

Skip List

A skip list for a set S of (key, value) pairs is a series of lists S_0, S_1, \dots, S_h such that

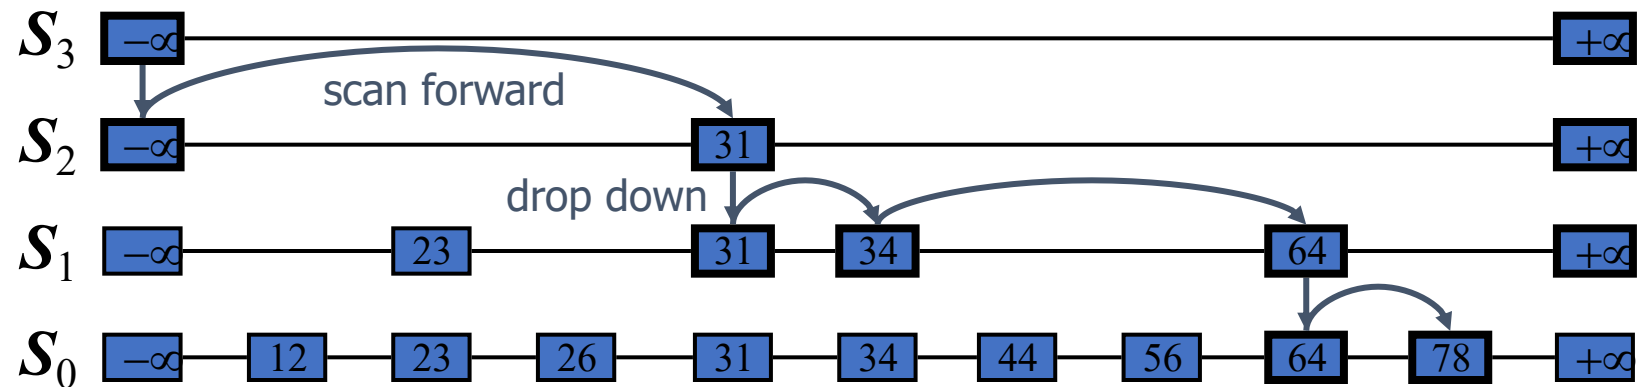
- each list contains special keys $-\infty$ and $+\infty$
- S_0 contains all keys of S in nondecreasing order
- Each list is a subsequence of the one before:
 $S_0 \supseteq S_1 \supseteq \dots \supseteq S_h$
- S_h only contains the two special keys



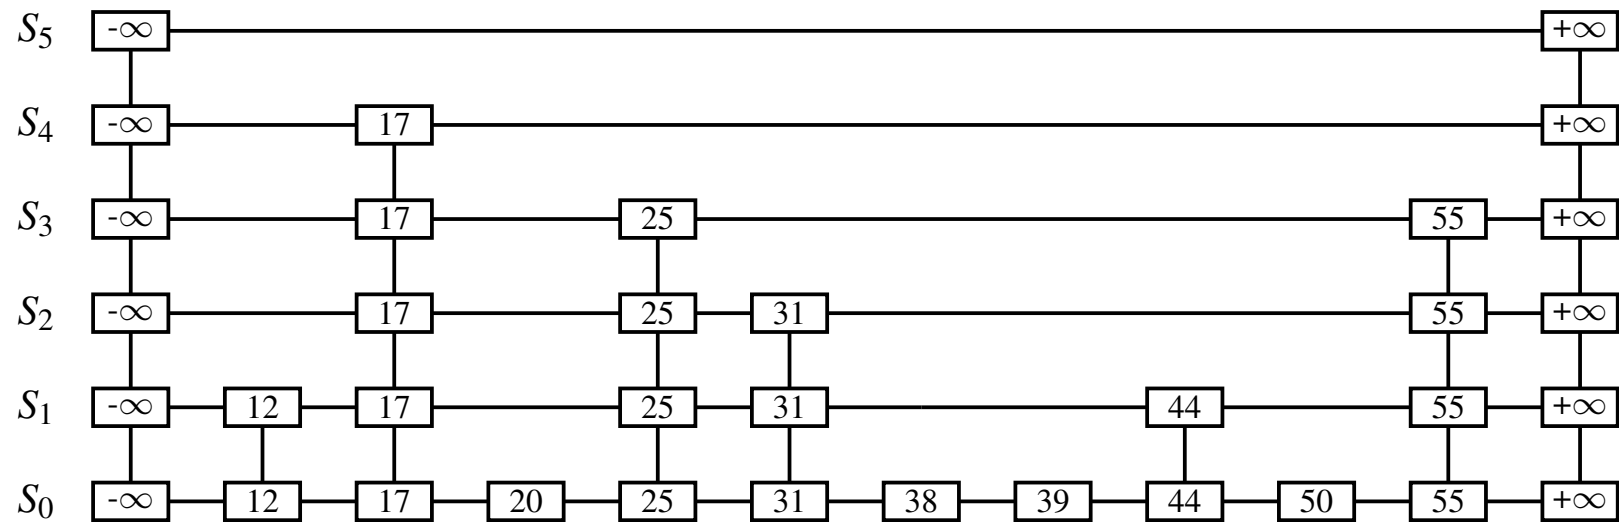
Search

Search for a key x in a skip list:


- start at the first position p of S_h
- compare x with $y = \text{next}(p)$
 - $x == y$ return y
 - $x > y$: scan forward $p = \text{next}(p)$
 - $x < y$: drop down $p = \text{below}(p)$



Example




Search for 50



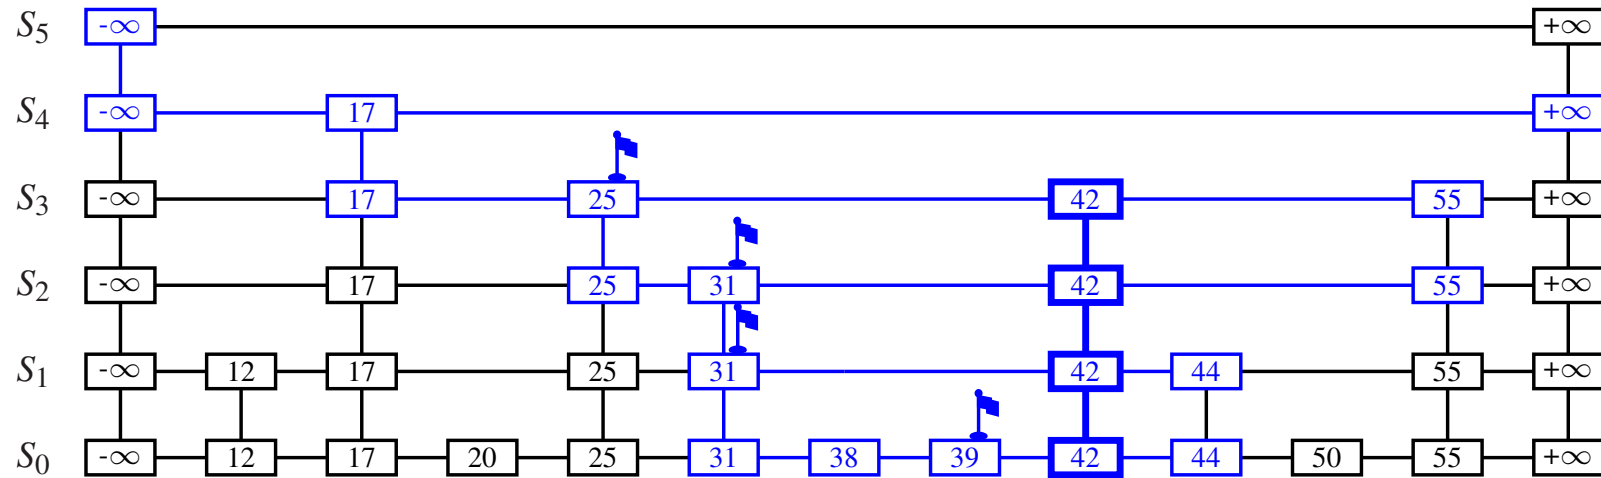
Skip List

- provides a clever compromise to realize a faster search on a sorted list (map)
- insertion is randomized
 - always insert into s_0
 - flip a coin for how many more lists to insert
 - expected runtime of $O(\log n)$
- search is $O(\log n)$ expected
- remove via search then up - $O(\log n)$ expected


Insert 42



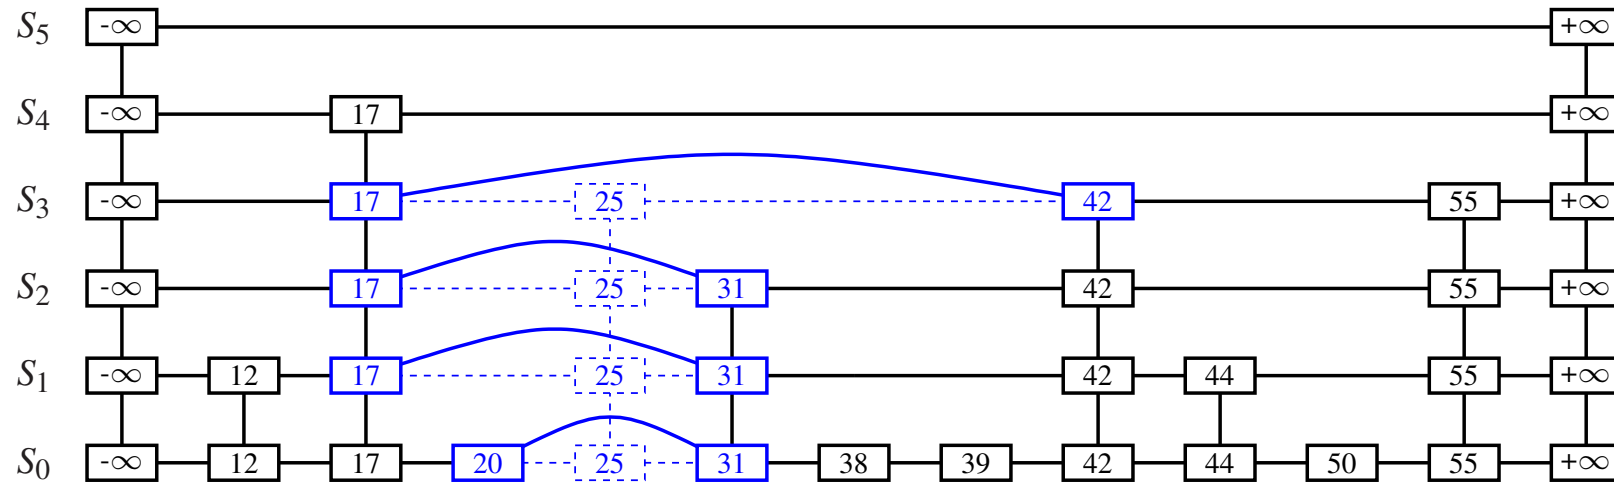
Insert 42



Remove 25



Remove 25



Skip List Analysis

Method	Running Time
size, isEmpty	$O(1)$
get	$O(\log n)$ expected
put	$O(\log n)$ expected
remove	$O(\log n)$ expected
firstEntry, lastEntry	$O(1)$
ceilingEntry, floorEntry lowerEntry, higherEntry	$O(\log n)$ expected
subMap	$O(s + \log n)$ expected, with s entries reported
entrySet, keySet, values	$O(n)$

Outline

Review Balanced Binary Trees

AVL Trees

Splay Trees

Skip List

Sets

Set

- A set is an unordered collection of elements, without duplicates
- A set supports an efficient search
- A hashtable is a set
- A multi-set (bag) allows duplicates
- A multi-map allows the same key to be mapped to multiple values

set ADT

`add(e)`: Adds the element e to S (if not already present).

`remove(e)`: Removes the element e from S (if it is present).

`contains(e)`: Returns whether e is an element of S .

`iterator()`: Returns an iterator of the elements of S .

There is also support for the traditional mathematical set operations of *union*, *intersection*, and *subtraction* of two sets S and T :

$$S \cup T = \{e: e \text{ is in } S \text{ or } e \text{ is in } T\},$$

$$S \cap T = \{e: e \text{ is in } S \text{ and } e \text{ is in } T\},$$

$$S - T = \{e: e \text{ is in } S \text{ and } e \text{ is not in } T\}.$$

`addAll(T)`: Updates S to also include all elements of set T , effectively replacing S by $S \cup T$.

`retainAll(T)`: Updates S so that it only keeps those elements that are also elements of set T , effectively replacing S by $S \cap T$.

`removeAll(T)`: Updates S by removing any of its elements that also occur in set T , effectively replacing S by $S - T$.

Implementation

- Recall that maps do not allow duplicate keys
- A set is simply a map in which keys have no associated values (or null)
- `java.util.HashSet`
- `java.util.Concurrent.ConcurrentSkipListSet`
- `java.util.TreeSet`

Java Built-ins: `java.util.*`

- **Linked List**
 - `LinkedList`
- **Stack**
 - `Stack` (linked)
- **Queue**
 - `ArrayDeque`
- **BST (unbalanced)**
 - none
- **Heap**
 - `PriorityQueue`
- **Hashtable**
 - `HashMap` (chained)
- **Set**
 - `HashSet`
- **Balanced BST**
 - `TreeMap` (R&B)
- **Search/Sort**
 - `Collections.binarySearch`
 - `Collections.sort`

Framework Diagram

